

```

import os
import json
import math
import random
import time
import re
import pathlib
import numpy as np
import pandas as pd
from PIL import Image
from pathlib import Path
from dataclasses import dataclass
import cv2

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
import torchvision.transforms as T
import timm

from sklearn.metrics import f1_score, accuracy_score, confusion_matrix
from sklearn.model_selection import GroupShuffleSplit # CRITICAL for patient split

from google.colab import drive

import zipfile
import os
from google.colab import drive

# 1. Mount Drive (if not already mounted)
drive.mount('/content/drive')

# 2. Define Paths
# Check which zip name you have in your Drive folder
zip_path_v1 = "/content/drive/MyDrive/ML Project/Clean_Split_Dataset.zip"
zip_path_v2 = "/content/drive/MyDrive/ML Project/Clean_Split_Dataset_Final.zip"

extract_to = "/content/via_dataset_fixed"

# 3. Check and Unzip
if os.path.exists(zip_path_v2):
    target_zip = zip_path_v2
    print(f"Found latest dataset: {target_zip}")
elif os.path.exists(zip_path_v1):
    target_zip = zip_path_v1
    print(f"Found dataset: {target_zip}")
else:
    target_zip = None
    print("ERROR: Could not find the processed dataset zip in your Drive.")
    print("Please re-run the 'Data Preparation' (Splitting & Augmentation) script.")

if target_zip:
    print(f"Unzipping to {extract_to}...")
    with zipfile.ZipFile(target_zip, 'r') as zip_ref:
        zip_ref.extractall(extract_to)
    print("Success! Data restored. You can now run the training script.")

# Verify the folder structure
if os.path.exists(os.path.join(extract_to, 'train')):
    print("Verified: 'train' folder exists.")
else:
    # Sometimes zip files have nested folders, let's fix that if it happens
    print("Checking for nested folders...")
    for root, dirs, files in os.walk(extract_to):
        if 'train' in dirs:
            nested_path = os.path.join(root, 'train')
            print(f"Found nested train folder at: {nested_path}")
            # You might need to update CFG.TRAIN_DIR in the training script to this path
            print(f"Update your CFG.TRAIN_DIR to: {nested_path}")

Mounted at /content/drive
Found latest dataset: /content/drive/MyDrive/ML Project/Clean_Split_Dataset_Final.zip
Unzipping to /content/via_dataset_fixed...
Success! Data restored. You can now run the training script.
Verified: 'train' folder exists.

# 1. Setup & Config
# -----
@dataclass
class CFG:
    # Path to the NEW split dataset you just created
    TRAIN_DIR = r"/content/via_dataset_fixed/train"
    TEST_DIR = r"/content/via_dataset_fixed/test"

    # Path to your original labels CSV (needed to map filename -> class)
    CSV_PATH = "/content/drive/MyDrive/ML Project/labels_combine.csv"

    NUM_CLASSES: int = 3
    IMAGE_SIZE: int = 384
    BATCH_SIZE: int = 32
    EPOCHS: int = 30
    LR: float = 1e-4
    WD: float = 1e-2
    SEED: int = 42
    MODEL_NAME: str = "efficientnet_b0"
    LOSS: str = "cross_entropy"
    OUT_DIR: str = "runs"

cfg = CFG()

```

```

def set_seed(seed=42):
    random.seed(seed); np.random.seed(seed)
    torch.manual_seed(seed); torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True

set_seed(cfg.SEED)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

Using device: cuda

# 2. Data Preparation Logic (THE BIG FIX)
# -----
def prepare_dataframes():
    """
    Scans the new folders, cleans filenames to find original keys,
    and maps them to labels from the master CSV.
    """

    # 1. Load Master Labels
    master_df = pd.read_csv(cfg.CSV_PATH)

    # Clean up column names/values if necessary
    # Assumes CSV has columns: 'File' (or image_path) and 'class'
    # We create a dictionary: { '1AFC1.jpg': 'Precancer', ... }

    # Normalize the filename column in CSV to be just the basename
    if 'image_path' in master_df.columns:
        master_df['filename_key'] = master_df['image_path'].apply(lambda x: os.path.basename(str(x)))
    elif 'File' in master_df.columns:
        master_df['filename_key'] = master_df['File'].apply(lambda x: os.path.basename(str(x)))
    else:
        raise ValueError("CSV must have 'image_path' or 'File' column")

    label_map = dict(zip(master_df['filename_key'], master_df['class']))

    # 2. Helper to scan a folder and build a DF
    def scan_folder(folder_path):
        data = []
        files = [f for f in os.listdir(folder_path) if f.lower().endswith('.jpg')]

        for f in files:
            # Logic: '1AFC1_rot.jpg' -> Original Key: '1AFC1.jpg'
            # Remove augmentation suffixes to find the label key
            base_name = f

            # Regex to strip _rot, _blur, _radial, etc.
            # Looks for the LAST occurrence of .jpg
            clean_name = re.sub(r'_\w+\.jpg', '', Path(f).stem)
            original_key = f"{clean_name}.jpg"

            # Parse Patient ID for splitting (e.g., 1AFC from 1AFC1.jpg)
            # Adjust regex based on your file format (Number+Letters or Letters+Number)
            # Trying robust search:
            match = re.search(r"([0-9]+[A-Z]+|[A-Z]+[0-9]+)", clean_name)
            patient_id = match.group(0) if match else "Unknown"

            if original_key in label_map:
                label = label_map[original_key]
                data.append({
                    "path": os.path.join(folder_path, f),
                    "label": label,
                    "patient_id": patient_id, # critical for GroupSplit
                    "is_aug": (f != original_key)
                })
            else:
                # Fallback: sometimes the key might be exact match
                pass

        return pd.DataFrame(data)

    print("Scanning Training Directory...")
    train_full_df = scan_folder(cfg.TRAIN_DIR)
    print(f"Found {len(train_full_df)} training images (including augmentations.)")

    print("Scanning Test Directory...")
    test_df = scan_folder(cfg.TEST_DIR)
    print(f"Found {len(test_df)} test images.")

    # 3. Map String Labels to Int
    mapping = {"Normal": 0, "Precancer": 1, "Cancer": 2}

    # Handle cases where CSV might already have numbers or different strings
    def map_label(x):
        if isinstance(x, int): return x
        return mapping.get(x, -1) # Returns -1 if not found

    train_full_df['class_num'] = train_full_df['label'].apply(map_label)
    test_df['class_num'] = test_df['label'].apply(map_label)

    # Filter out bad labels
    train_full_df = train_full_df[train_full_df['class_num'] != -1]
    test_df = test_df[test_df['class_num'] != -1]

    # 4. Patient-Aware Validation Split
    # We split the TRAIN folder into Train/Val, but strictly by PATIENT ID
    # This prevents '1AFC1_rot.jpg' being in Train and '1AFC1.jpg' in Val
    gss = GroupShuffleSplit(n_splits=1, test_size=0.15, random_state=cfg.SEED)

    train_idx, val_idx = next(gss.split(train_full_df, groups=train_full_df['patient_id']))

    train_df = train_full_df.iloc[train_idx].reset_index(drop=True)
    val_df = train_full_df.iloc[val_idx].reset_index(drop=True)

```

```

print("-" * 30)
print(f"Final Train Size: {len(train_df)} (Augmented)")
print(f"Final Val Size: {len(val_df)} (Augmented subset)")
print(f"Final Test Size: {len(test_df)} (Original only)")
print("-" * 30)

return train_df, val_df, test_df

# 3. Dataset Class
# -----
clip_limit = 13
class CervicalDataset(Dataset):
    def __init__(self, df, transform=None):
        self.df = df
        self.transform = transform

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = row['class_num']

        img = Image.open(img_path).convert("RGB")

        if self.transform:
            img = self.transform(img)

        return img, torch.tensor(label, dtype=torch.long)

# 4. Transforms
# Define a custom CLAHE transform class
class ApplyCLAHE(object):
    def __init__(self, clip_limit=clip_limit, tile_grid_size=(8, 8)):
        self.clahe = cv2.createCLAHE(clipLimit=clip_limit, tileGridSize=tile_grid_size)

    def __call__(self, img):
        # Convert PIL to Numpy
        img_np = np.array(img)

        # Convert RGB to LAB color space (L = Lightness)
        lab = cv2.cvtColor(img_np, cv2.COLOR_RGB2LAB)
        l, a, b = cv2.split(lab)

        # Apply CLAHE to the L-channel only (contrast)
        cl = self.clahe.apply(l)

        # Merge back
        limg = cv2.merge((cl, a, b))

        # Convert back to RGB
        final = cv2.cvtColor(limg, cv2.COLOR_LAB2RGB)
        return Image.fromarray(final)

    # Update your get_transforms
    def get_transforms(image_size):
        train_tfms = T.Compose([
            ApplyCLAHE(clip_limit=clip_limit), # <--- ADD THIS FIRST
            T.Resize((image_size, image_size)),
            T.RandomHorizontalFlip(p=0.5),
            T.RandomVerticalFlip(p=0.5), # Add vertical flip too (anatomy has no "up")
            T.ToTensor(),
            T.Normalize(mean=(0.485, 0.456, 0.406), std=(0.229, 0.224, 0.225)),
        ])
        # Apply to Validation too so the model sees the same "enhanced" view
        val_tfms = T.Compose([
            ApplyCLAHE(clip_limit=clip_limit), # <--- ADD THIS
            T.Resize((image_size, image_size)),
            T.ToTensor(),
            T.Normalize(mean=(0.485, 0.456, 0.406), std=(0.229, 0.224, 0.225)),
        ])
        return train_tfms, val_tfms

# 5. Training Loop Components
# -----
def train_one_epoch(model, loader, criterion, optimizer):
    model.train()
    losses, accs = [], []
    for x, y in loader:
        x, y = x.to(device), y.to(device)
        optimizer.zero_grad()
        logits = model(x)
        loss = criterion(logits, y)
        loss.backward()
        optimizer.step()

        preds = logits.argmax(1)
        acc = (preds == y).float().mean()
        losses.append(loss.item())
        accs.append(acc.item())
    return np.mean(losses), np.mean(accs)

@torch.no_grad()
def validate(model, loader, criterion):
    model.eval()
    losses, accs, all_preds, all_targets = [], [], [], []
    for x, y in loader:
        x, y = x.to(device), y.to(device)
        logits = model(x)
        loss = criterion(logits, y)

```

```

preds = logits.argmax(1)

losses.append(loss.item())
accs.append((preds == y).float().mean().item())
all_preds.extend(preds.cpu().numpy())
all_targets.extend(y.cpu().numpy())

f1 = f1_score(all_targets, all_preds, average='macro')
return np.mean(losses), np.mean(accs), f1, all_preds, all_targets

# 6. Main Execution
# -----
def main():
    # A. Prepare Data
    train_df, val_df, test_df = prepare_dataframes()

    train_tfms, val_tfms = get_transforms(cfg.IMAGE_SIZE)

    train_ds = CervicalDataset(train_df, transform=train_tfms)
    val_ds = CervicalDataset(val_df, transform=val_tfms)
    test_ds = CervicalDataset(test_df, transform=val_tfms)

    # B. Weighted Sampler for Imbalance
    # Calculate weights based on Class counts in Training set
    class_counts = train_df['class_num'].value_counts().sort_index()
    print("Class Counts:", class_counts.to_dict())

    sample_weights = [1.0 / class_counts[label] for label in train_df['class_num']]
    sampler = WeightedRandomSampler(sample_weights, num_samples=len(sample_weights), replacement=True)

    train_loader = DataLoader(train_ds, batch_size=cfg.BATCH_SIZE, sampler=sampler, num_workers=2)
    val_loader = DataLoader(val_ds, batch_size=cfg.BATCH_SIZE, shuffle=False, num_workers=2)
    test_loader = DataLoader(test_ds, batch_size=cfg.BATCH_SIZE, shuffle=False, num_workers=2)

    # C. Model Setup (3-CLASS, PARTIAL UNFREEZE)
    # -----

    # 1. Create Model
    model = timm.create_model(cfg.MODEL_NAME, pretrained=True, num_classes=3) # Force 3 classes
    model = model.to(device)

    # 2. Intelligent Unfreezing (The "Goldilocks" Fix)
    # We freeze the first 60% (Generic shapes) and train the last 40% (Medical textures)

    # First, freeze everything
    for param in model.parameters():
        param.requires_grad = False

    # Unfreeze the Head (Classifier)
    for param in model.classifier.parameters():
        param.requires_grad = True

    # Unfreeze the Last 3 Blocks of EfficientNet
    # These are the layers that see "patterns" and "textures"
    # (EfficientNet B0 has 7 blocks total)
    for param in model.blocks[-1].parameters(): # Block 6
        param.requires_grad = True
    for param in model.blocks[-2].parameters(): # Block 5
        param.requires_grad = True
    for param in model.blocks[-3].parameters(): # Block 4
        param.requires_grad = True

    # Check what we are training
    trainable = sum(p.numel() for p in model.parameters() if p.requires_grad)
    total = sum(p.numel() for p in model.parameters())
    print(f"Partial Unfreeze: Training {trainable:,} of {total:,} params.")

    # 3. Loss Function Upgrade: Label Smoothing
    # Standard CrossEntropy forces the model to shout "100% CANCER!"
    # Label Smoothing tells it to say "90% Cancer, 10% Other", which reduces overfitting.
    from timm.loss import LabelSmoothingCrossEntropy

    # We still use class weights because of imbalance
    counts = train_df['class_num'].value_counts().sort_index()
    weights = torch.tensor([1.0/counts[0], 1.0/counts[1], 1.0/counts[2]], dtype=torch.float32).to(device)
    weights = weights / weights.sum()

    # Note: LabelSmoothing in timm usually doesn't take class weights directly,
    # so we often use standard CrossEntropy with label_smoothing arg if using PyTorch 1.10+,
    # OR we use the weighted approach. Let's stick to Weighted CE for now but add strong regularization.
    criterion = nn.CrossEntropyLoss(weight=weights, label_smoothing=0.1)

    # 4. Optimizer
    # We use a smaller Learning Rate because we are training body layers now
    optimizer = torch.optim.AdamW(model.parameters(), lr=5e-4, weight_decay=0.02)
    scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=cfg.EPOCHS)

    # D. Training Loop
    out_dir = Path(cfg.OUT_DIR) / f"exp_{time.strftime('%m%d_%H%M')}"
    out_dir.mkdir(parents=True, exist_ok=True)

    best_f1 = -1.0
    history = {'train_loss': [], 'val_loss': [], 'val_f1': []}

    print("\nStarting Training...")
    for epoch in range(1, cfg.EPOCHS + 1):
        t_loss, t_acc = train_one_epoch(model, train_loader, criterion, optimizer)
        v_loss, v_acc, v_f1, _, _ = validate(model, val_loader, criterion)

        scheduler.step()

        history['train_loss'].append(t_loss)

```

```

history['val_loss'].append(v_loss)
history['val_f1'].append(v_f1)

print(f"Epoch {epoch:02d} | Train Loss: {t_loss:.4f} Acc: {t_acc:.3f} | Val Loss: {v_loss:.4f} F1: {v_f1:.3f}")

if v_f1 > best_f1:
    best_f1 = v_f1
    torch.save(model.state_dict(), out_dir / "best_model.pth")

print("\nTraining Complete. Best Val F1: {best_f1:.3f}")

# E. Final Test on Pure Unseen Data
print("\nRunning Final Test on Clean Test Set...")
model.load_state_dict(torch.load(out_dir / "best_model.pth"))
test_loss, test_acc, test_f1, test_preds, test_targets = validate(model, test_loader, criterion)

print(f"TEST RESULTS -> Acc: {test_acc:.3f} | F1: {test_f1:.3f}")

# Save History
json.dump(history, open(out_dir / "history.json", "w"))

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

# 1. Define the mapping explicitly
mapping = {"Normal": 0, "Precancer": 1, "Cancer": 2}
class_names = list(mapping.keys())

# 2. Plot
cm = confusion_matrix(test_targets, test_preds)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f"Confusion Matrix (Test Acc: {test_acc:.3f})")
plt.show()

# 3. Print a report to see which class is killing you
from sklearn.metrics import classification_report
print(classification_report(test_targets, test_preds, target_names=class_names))

if __name__ == "__main__":
    main()

```

Scanning Training Directory...

Found 745 training images (including augmentations).

Scanning Test Directory...

Found 37 test images.

Final Train Size: 630 (Augmented)

Final Val Size: 115 (Augmented subset)

Final Test Size: 37 (Original only)

Class Counts: {0: 345, 1: 210, 2: 75}

Partial Unfreeze: Training 3,290,571 of 4,011,391 params.

Starting Training...

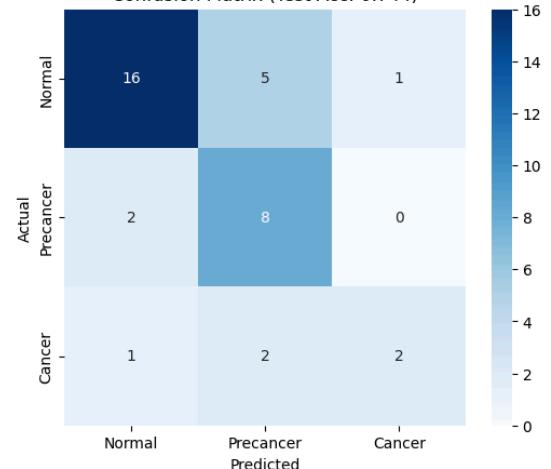
Epoch	Train Loss	Acc	Val Loss	F1
01	1.1642	0.683	2.3777	0.416
02	0.6384	0.843	1.7316	0.478
03	0.5701	0.870	1.5730	0.504
04	0.4225	0.937	1.4518	0.425
05	0.4083	0.930	1.3675	0.494
06	0.3378	0.964	1.3815	0.535
07	0.3368	0.959	1.2807	0.529
08	0.3328	0.970	1.3198	0.499
09	0.2914	0.976	1.3226	0.463
10	0.3037	0.976	1.2746	0.506
11	0.2990	0.978	1.2692	0.502
12	0.2919	0.984	1.2527	0.474
13	0.2819	0.989	1.2686	0.481
14	0.2917	0.989	1.2639	0.493
15	0.3054	0.983	1.2490	0.489

Training Complete. Best Val F1: 0.535

Running Final Test on Clean Test Set...

TEST RESULTS -> Acc: 0.744 | F1: 0.640

Confusion Matrix (Test Acc: 0.744)



precision recall f1-score support

	precision	recall	f1-score	support
Normal	0.84	0.73	0.78	22
Precancer	0.53	0.80	0.64	10
Cancer	0.67	0.40	0.50	5
accuracy			0.70	37
macro avg	0.68	0.64	0.64	37
weighted avg	0.73	0.70	0.70	37

```
# 6. Main Execution
# -----
def main():
    # A. Prepare Data
    train_df, val_df, test_df = prepare_dataframes()

    train_tfms, val_tfms = get_transforms(cfg.IMAGE_SIZE)

    train_ds = CervicalDataset(train_df, transform=train_tfms)
    val_ds = CervicalDataset(val_df, transform=val_tfms)
    test_ds = CervicalDataset(test_df, transform=val_tfms)

    # B. Weighted Sampler for Imbalance
    # Calculate weights based on Class counts in Training set
    class_counts = train_df['class_num'].value_counts().sort_index()
    print("Class Counts:", class_counts.to_dict())

    sample_weights = [1.0 / class_counts[label] for label in train_df['class_num']]
    sampler = WeightedRandomSampler(sample_weights, num_samples=len(sample_weights), replacement=True)

    train_loader = DataLoader(train_ds, batch_size=cfg.BATCH_SIZE, sampler=sampler, num_workers=2)
    val_loader = DataLoader(val_ds, batch_size=cfg.BATCH_SIZE, shuffle=False, num_workers=2)
    test_loader = DataLoader(test_ds, batch_size=cfg.BATCH_SIZE, shuffle=False, num_workers=2)

    # C. Model Setup (3-CLASS, PARTIAL UNFREEZE)
    # -----

    # 1. Create Model
    model = timm.create_model(cfg.MODEL_NAME, pretrained=True, num_classes=3) # Force 3 classes
    model = model.to(device)

    # 2. Intelligent Unfreezing (The "Goldilocks" Fix)
    # We freeze the first 60% (Generic shapes) and train the last 40% (Medical textures)

    # First, freeze everything
    for param in model.parameters():
        param.requires_grad = False
```

```

# Unfreeze the Head (Classifier)
for param in model.classifier.parameters():
    param.requires_grad = True

# Unfreeze the Last 3 Blocks of EfficientNet
# These are the layers that see "patterns" and "textures"
# (EfficientNet B0 has 7 blocks total)
for param in model.blocks[-1].parameters(): # Block 6
    param.requires_grad = True
for param in model.blocks[-2].parameters(): # Block 5
    param.requires_grad = True
for param in model.blocks[-3].parameters(): # Block 4
    param.requires_grad = True

# Check what we are training
trainable = sum(p.numel() for p in model.parameters() if p.requires_grad)
total = sum(p.numel() for p in model.parameters())
print(f"Partial Unfreeze: Training {trainable:,} of {total:,} params.")

# 3. Loss Function Upgrade: Label Smoothing
# Standard CrossEntropy forces the model to shout "100% CANCER!"
# Label Smoothing tells it to say "90% Cancer, 10% Other", which reduces overfitting.
from timm.loss import LabelSmoothingCrossEntropy

# We still use class weights because of imbalance
counts = train_df['class_num'].value_counts().sort_index()
weights = torch.tensor([1.0/counts[0], 1.0/counts[1], 1.0/counts[2]], dtype=torch.float32).to(device)
weights = weights / weights.sum()

# Note: LabelSmoothing in timm usually doesn't take class weights directly,
# so we often use standard CrossEntropy with label_smoothing arg if using PyTorch 1.10+,
# OR we use the weighted approach. Let's stick to Weighted CE for now but add strong regularization.
criterion = nn.CrossEntropyLoss(weight=weights, label_smoothing=0.1)

# 4. Optimizer
# We use a smaller Learning Rate because we are training body layers now
optimizer = torch.optim.AdamW(model.parameters(), lr=5e-4, weight_decay=0.02)
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=cfg.EPOCHS)

# D. Training Loop
out_dir = Path(cfg.OUT_DIR) / f"exp_{time.strftime('%m%d_%H%M')}"
out_dir.mkdir(parents=True, exist_ok=True)

best_f1 = -1.0
history = {'train_loss': [], 'val_loss': [], 'val_f1': []}

print("\nStarting Training...")
for epoch in range(1, cfg.EPOCHS + 1):
    t_loss, t_acc = train_one_epoch(model, train_loader, criterion, optimizer)
    v_loss, v_acc, v_f1, _, _ = validate(model, val_loader, criterion)

    scheduler.step()

    history['train_loss'].append(t_loss)
    history['val_loss'].append(v_loss)
    history['val_f1'].append(v_f1)

    print(f"Epoch {epoch:02d} | Train Loss: {t_loss:.4f} Acc: {t_acc:.3f} | Val Loss: {v_loss:.4f} F1: {v_f1:.3f}")

    if v_f1 > best_f1:
        best_f1 = v_f1
        torch.save(model.state_dict(), out_dir / "best_model.pth")

print("\nTraining Complete. Best Val F1: {best_f1:.3f}")

# E. Final Test on Pure Unseen Data
print("\nRunning Final Test on Clean Test Set...")
model.load_state_dict(torch.load(out_dir / "best_model.pth"))
test_loss, test_acc, test_f1, test_preds, test_targets = validate(model, test_loader, criterion)

print(f"\nTEST RESULTS -> Acc: {test_acc:.3f} | F1: {test_f1:.3f}")

# Save History
json.dump(history, open(out_dir / "history.json", "w"))

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

# 1. Define the mapping explicitly
mapping = {"Normal": 0, "Precancer": 1, "Cancer": 2}
class_names = list(mapping.keys())

# 2. Plot
cm = confusion_matrix(test_targets, test_preds)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f"Confusion Matrix (Test Acc: {test_acc:.3f})")
plt.show()

# 3. Print a report to see which class is killing you
from sklearn.metrics import classification_report
print(classification_report(test_targets, test_preds, target_names=class_names))

if __name__ == "__main__":
    main()

```

```
Scanning Training Directory...
Found 745 training images (including augmentations).
Scanning Test Directory...
Found 37 test images.
```

```
Final Train Size: 630 (Augmented)
Final Val Size: 115 (Augmented subset)
Final Test Size: 37 (Original only)
```

```
Class Counts: {0: 345, 1: 210, 2: 75}
Partial Unfreeze: Training 3,290,571 of 4,011,391 params.
```

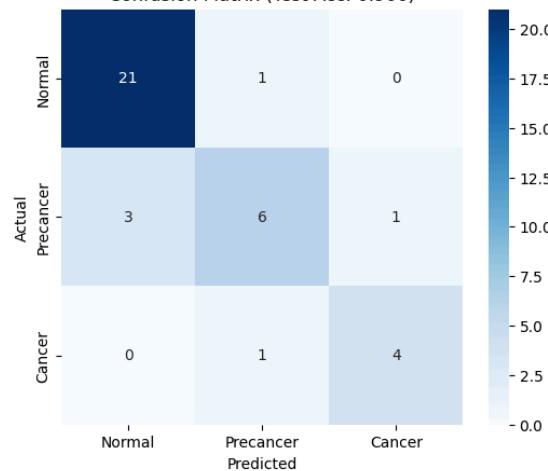
Starting Training...

Epoch	Train Loss	Acc	Val Loss	F1
01	1.1548	0.688	2.0364	0.553
02	0.6745	0.839	1.7209	0.482
03	0.4976	0.869	1.4505	0.532
04	0.4376	0.912	1.4984	0.490
05	0.3999	0.949	1.2140	0.516
06	0.3436	0.954	1.0650	0.594
07	0.3720	0.960	1.0713	0.573
08	0.3210	0.980	1.1029	0.568
09	0.3332	0.976	1.0555	0.529
10	0.2848	0.984	1.0700	0.539
11	0.3015	0.990	1.0500	0.601
12	0.2785	0.995	1.0166	0.620
13	0.2834	0.995	1.0089	0.586
14	0.2744	0.991	0.9964	0.551
15	0.2847	0.997	1.0341	0.591
16	0.2851	0.995	1.0562	0.576
17	0.2810	0.996	1.0338	0.588
18	0.2721	0.998	1.0252	0.566
19	0.2657	0.998	1.0296	0.632
20	0.2806	1.000	1.0414	0.581
21	0.2788	1.000	1.0236	0.560
22	0.2678	0.998	1.0263	0.565
23	0.2791	1.000	1.0216	0.569
24	0.2738	1.000	1.0290	0.588
25	0.2609	0.998	1.0353	0.554
26	0.2746	0.995	1.0178	0.562
27	0.2648	1.000	1.0139	0.551
28	0.2499	1.000	1.0155	0.618
29	0.2560	0.998	1.0206	0.586
30	0.2609	1.000	1.0269	0.607

Training Complete. Best Val F1: 0.632

Running Final Test on Clean Test Set...
TEST RESULTS -> Acc: 0.906 | F1: 0.793

Confusion Matrix (Test Acc: 0.906)



	precision	recall	f1-score	support
Normal	0.88	0.95	0.91	22
Precancer	0.75	0.60	0.67	10
Cancer	0.80	0.80	0.80	5
accuracy			0.84	37
macro avg	0.81	0.78	0.79	37
weighted avg	0.83	0.84	0.83	37

```
# # 6. Main Execution
# #
# def main():
#     # A. Prepare Data
#     train_df, val_df, test_df = prepare_dataframes()

#     train_tfms, val_tfms = get_transforms(cfg.IMAGE_SIZE)

#     train_ds = CervicalDataset(train_df, transform=train_tfms)
#     val_ds = CervicalDataset(val_df, transform=val_tfms)
#     test_ds = CervicalDataset(test_df, transform=val_tfms)

#     # B. Weighted Sampler for Imbalance
#     # Calculate weights based on Class counts in Training set
#     class_counts = train_df['class_num'].value_counts().sort_index()
#     print("Class Counts:", class_counts.to_dict())

#     sample_weights = [1.0 / class_counts[label] for label in train_df['class_num']]
#     sampler = WeightedRandomSampler(sample_weights, num_samples=len(sample_weights), replacement=True)

#     train_loader = DataLoader(train_ds, batch_size=cfg.BATCH_SIZE, sampler=sampler, num_workers=2)
#     val_loader = DataLoader(val_ds, batch_size=cfg.BATCH_SIZE, shuffle=False, num_workers=2)
#     test_loader = DataLoader(test_ds, batch_size=cfg.BATCH_SIZE, shuffle=False, num_workers=2)
```

```

# # C. Model Setup (3-CLASS, PARTIAL UNFREEZE)
# -----
#
# # 1. Create Model
# model = timm.create_model(cfg.MODEL_NAME, pretrained=True, num_classes=3) # Force 3 classes
# model = model.to(device)

# # 2. Intelligent Unfreezing (The "Goldilocks" Fix)
# # We freeze the first 60% (Generic shapes) and train the last 40% (Medical textures)

# # First, freeze everything
# for param in model.parameters():
#     param.requires_grad = False

# # Unfreeze the Head (Classifier)
# for param in model.classifier.parameters():
#     param.requires_grad = True

# # Unfreeze the Last 3 Blocks of EfficientNet
# # These are the layers that see "patterns" and "textures"
# # (EfficientNet B0 has 7 blocks total)
# for param in model.blocks[-1].parameters(): # Block 6
#     param.requires_grad = True
# for param in model.blocks[-2].parameters(): # Block 5
#     param.requires_grad = True
# # for param in model.blocks[-3].parameters(): # Block 4
#     param.requires_grad = True

# # Check what we are training
# trainable = sum(p.numel() for p in model.parameters() if p.requires_grad)
# total = sum(p.numel() for p in model.parameters())
# print(f"Partial Unfreeze: Training {trainable:,} of {total:,} params.")

# # 3. Loss Function Upgrade: Label Smoothing
# # Standard CrossEntropy forces the model to shout "100% CANCER!"
# # Label Smoothing tells it to say "90% Cancer, 10% Other", which reduces overfitting.
# from timm.loss import LabelSmoothingCrossEntropy

# # We still use class weights because of imbalance
# counts = train_df['class_num'].value_counts().sort_index()
# weights = torch.tensor([1.0/counts[0], 1.0/counts[1], 1.0/counts[2]], dtype=torch.float32).to(device)
# weights = weights / weights.sum()

# # Note: LabelSmoothing in timm usually doesn't take class weights directly,
# # so we often use standard CrossEntropy with label_smoothing arg if using PyTorch 1.10+,
# # OR we use the weighted approach. Let's stick to Weighted CE for now but add strong regularization.
# criterion = nn.CrossEntropyLoss(weight=weights, label_smoothing=0.1)

# # 4. Optimizer
# # We use a smaller Learning Rate because we are training body layers now
# optimizer = torch.optim.AdamW(model.parameters(), lr=5e-4, weight_decay=0.02)
# scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=cfg.EPOCHS)

# # D. Training Loop
# out_dir = Path(cfg.OUT_DIR) / f'exp_{time.strftime("%m%d_%H%M")}'
# out_dir.mkdir(parents=True, exist_ok=True)

# best_f1 = -1.0
# history = {'train_loss': [], 'val_loss': [], 'val_f1': []}

# print("\nStarting Training...")
# for epoch in range(1, cfg.EPOCHS + 1):
#     t_loss, t_acc = train_one_epoch(model, train_loader, criterion, optimizer)
#     v_loss, v_acc, v_f1, _, _ = validate(model, val_loader, criterion)

#     scheduler.step()

#     history['train_loss'].append(t_loss)
#     history['val_loss'].append(v_loss)
#     history['val_f1'].append(v_f1)

#     print(f"Epoch {epoch:02d} | Train Loss: {t_loss:.4f} Acc: {t_acc:.3f} | Val Loss: {v_loss:.4f} F1: {v_f1:.3f}")

#     if v_f1 > best_f1:
#         best_f1 = v_f1
#         torch.save(model.state_dict(), out_dir / "best_model.pth")

# print(f"\nTraining Complete. Best Val F1: {best_f1:.3f}")

# # E. Final Test on Pure Unseen Data
# print("\nRunning Final Test on Clean Test Set...")
# model.load_state_dict(torch.load(out_dir / "best_model.pth"))
# test_loss, test_acc, test_f1, test_preds, test_targets = validate(model, test_loader, criterion)

# print(f"TEST RESULTS -> Acc: {test_acc:.3f} | F1: {test_f1:.3f}")

# # Save History
# json.dump(history, open(out_dir / "history.json", "w"))

# import matplotlib.pyplot as plt
# import seaborn as sns
# from sklearn.metrics import confusion_matrix

# # 1. Define the mapping explicitly
# mapping = {"Normal": 0, "Precancer": 1, "Cancer": 2}
# class_names = list(mapping.keys())

# # 2. Plot
# cm = confusion_matrix(test_targets, test_preds)
# plt.figure(figsize=(6,5))
# sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
#             xticklabels=class_names, yticklabels=class_names)
# plt.xlabel('Predicted')

```

```
# plt.xlabel('Predicted')
# plt.ylabel('Actual')
# plt.title(f"Confusion Matrix (Test Acc: {test_acc:.3f})")
# plt.show()

# # 3. Print a report to see which class is killing you
# from sklearn.metrics import classification_report
# print(classification_report(test_targets, test_preds, target_names=class_names))

# # ... previous code (confusion matrix plotting) ...

# # --- NEW: SAVE AS PKL ---
# print("\nSaving model as .pkl (Pickle)...")
# pkl_path = out_dir / "final_model_full.pkl"

# # We save the entire model object, not just the weights
# torch.save(model, pkl_path)

# print(f"SUCCESS! Complete model saved to: {pkl_path}")
# print(f"Best Weights (.pth) are at: {out_dir / 'best_model.pth'}")

# # Optional: Copy to Drive immediately so you don't lose it
# import shutil
# drive_save_path = Path("/content/drive/MyDrive/ML Project/final_model_full.pkl")
# shutil.copy(pkl_path, drive_save_path)
# print(f"Backup saved to Google Drive: {drive_save_path}")
```